Linux Device Drivers, 2nd Edition

Chapter 14
Network Drivers

Contents:

We are now through discussing char and block drivers and are ready to move on to the fascinating world of networking. Network interfaces are the third standard class of Linux devices, and this chapter describes how they interact with the rest of the kernel.

The role of a network interface within the system is similar to that of a mounted block device. A block device registers its features in the blk_dev array and other kernel structures, and it then "transmits" and "receives" blocks on request, by means of its request function. Similarly, a network interface must register itself in specific data structures in order to be invoked when packets are exchanged with the outside world.

There are a few important differences between mounted disks and packet-delivery interfaces. To begin with, a disk exists as a special file in the /dev directory, whereas a network interface has no such entry point.

The normal file operations (read, write, and so on) do not make sense when applied to network interfaces, so it is not possible to apply the Unix "everything is a file" approach to them. Thus, network interfaces exist in their own namespace and export a different set of operations.

Although you may object that applications use the read and write system calls when using sockets, those calls act on a software object that is distinct from the interface. Several hundred sockets can be multiplexed on the same physical interface.

But the most important difference between the two is that block drivers operate only in response to requests from the kernel, whereas network drivers receive packets asynchronously from the outside. Thus, while a block driver is asked to send a buffer toward the kernel, the network device asksto push incoming packets toward the kernel. The kernel interface for network drivers is designed for this different mode of operation.

Network drivers also have to be prepared to support a number of administrative tasks, such as setting addresses, modifying transmission parameters, and maintaining traffic and error statistics. The API for network drivers reflects this need, and thus looks somewhat different from the interfaces we have seen so far.

The network subsystem of the Linux kernel is designed to be completely protocol independent. This applies to both networking protocols (IP versus IPX or other protocols) and hardware protocols (Ethernet versus token ring, etc.). Interaction between a network driver and the kernel proper deals with one network packet at a time; this allows protocol issues to be hidden neatly from the driver and the physical transmission to be hidden from the protocol.

This chapter describes how the network interfaces fit in with the rest of the Linux kernel and shows a memory-based modularized network interface, which is called (you guessed it) snull. To simplify the discussion, the interface uses the Ethernet hardware protocol and transmits IP packets. The knowledge you acquire from examining snull can be readily applied to protocols other than IP, and writing a non-Ethernet driver is only different in tiny details related to the actual network protocol.

This chapter doesn't talk about IP numbering schemes, network protocols, or other general networking concepts. Such topics are not (usually) of concern to the driver writer, and it's impossible to offer a satisfactory overview of networking technology in less than a few hundred pages. The interested reader is urged to refer to other books describing networking issues.

The networking subsystem has seen many changes over the years as the kernel developers have striven to provide the best performance possible. The bulk of this chapter describes network drivers as they are implemented in the 2.4 kernel. Once again, the sample code works on the 2.0 and 2.2 kernels as well, and we cover the differences between those kernels and 2.4 at the end of the chapter.

One note on terminology is called for before getting into network devices. The networking world uses the term octet to refer to a group of eight bits, which is generally the smallest unit understood by networking devices and protocols. The term byte is almost never encountered in this context. In keeping with standard usage, we will use octet when talking about networking devices.

How snull Is Designed

This section discusses the design concepts that led to the snull network interface. Although this information might appear to be of marginal use, failing to understand this driver might lead to problems while playing with the sample code.

The first, and most important, design decision was that the sample interfaces should remain independent of real hardware, just like most of the sample code used in this book. This constraint led to something that resembles the loopback interface. snull is not a loopback interface, however; it simulates conversations with real remote hosts in order to better demonstrate the task of writing a network driver. The Linux loopback driver is actually quite simple; it can be found in drivers/net/loopback.c.

Another feature of snull is that it supports only IP traffic. This is a consequence of the internal workings of the interface -- snull has to look inside and interpret the packets to properly emulate a pair of hardware interfaces. Real interfaces don't depend on the protocol being transmitted, and this limitation of snull doesn't affect the fragments of code that are shown in this chapter.

Assigning IP Numbers

The snull module creates two interfaces. These interfaces are different from a simple loopback in that whatever you transmit through one of the interfaces loops back to the other one, not to itself. It looks like you have two external links, but actually your computer is replying to itself.

Unfortunately, this effect can't be accomplished through IP-number assignment alone, because the kernel wouldn't send out a packet through interface A that was directed to its own interface B. Instead, it would use the loopback channel without passing through snull. To be able to establish a communication through the snull interfaces, the source and destination addresses need to be modified during data transmission. In other words, packets sent through one of the interfaces should be received by the other, but the receiver of the outgoing packet shouldn't be recognized as the local host. The same applies to the source address of received packets.

To achieve this kind of "hidden loopback," the snull interface toggles the least significant bit of the third octet of both the source and destination addresses; that is, it changes both the network number and the host number of class C IP numbers. The net effect is that packets

sent to network A (connected to sn0, the first interface) appear on the sn1 interface as packets belonging to network B.

To avoid dealing with too many numbers, let's assign symbolic names to the IP numbers involved:

- snullnet0 is the class C network that is connected to the sn0 interface. Similarly, snullnet1 is the network connected to sn1. The addresses of these networks should differ only in the least significant bit of the third octet.

- local0 is the IP address assigned to the sn0 interface; it belongs to snullnet0. The address associated with sn1 is local1. local0 and local1 must differ in the least significant bit of their third octet and in the fourth octet.

- remote0 is a host in snullnet0, and its fourth octet is the same as that of local1. Any packet sent to remote0 will reach local1 after its class C address has been modified by the interface code. The host remote1 belongs to snullnet1, and its fourth octet is the same as that of local0.

The operation of the snull interfaces is depicted in Figure 14-1, in which the hostname associated with each interface is printed near the interface name.

Figure 14-1. How a host sees its interfaces

Here are possible values for the network numbers. Once you put these lines in /etc/networks, you can call your networks by name. The values shown were chosen from the range of numbers reserved for private use.

```
snullnet0        192.168.0.0
snullnet1        192.168.1.0
```

The following are possible host numbers to put into /etc/hosts:

```
192.168.0.1    local0
```

```
192.168.0.2   remote0
192.168.1.2   local1
192.168.1.1   remote1
```

The important feature of these numbers is that the host portion of `local0` is the same as that of `remote1`, and the host portion of `local1` is the same as that of `remote0`. You can use completely different numbers as long as this relationship applies.

Be careful, however, if your computer is already connected to a network. The numbers you choose might be real Internet or intranet numbers, and assigning them to your interfaces will prevent communication with the real hosts. For example, although the numbers just shown are not routable Internet numbers, they could already be used by your private network if it lives behind a firewall.

Whatever numbers you choose, you can correctly set up the interfaces for operation by issuing the following commands:

```
ifconfig sn0 local0
ifconfig sn1 local1
case "`uname -r`" in 2.0.*)
    route add -net snullnet0 dev sn0
    route add -net snullnet1 dev sn1
esac
```

There is no need to invoke route with 2.2 and later kernels because the route is automatically added. Also, you may need to add the `netmask` `255.255.255.0` parameter if the address range chosen is not a class C range.

At this point, the "remote" end of the interface can be reached. The following screendump shows how a host reaches `remote0` and `remote1` through the snull interface.

```
morgana% ping -c 2 remote0
64 bytes from 192.168.0.99: icmp_seq=0 ttl=64 time=1.6 ms
64 bytes from 192.168.0.99: icmp_seq=1 ttl=64 time=0.9 ms
2 packets transmitted, 2 packets received, 0% packet loss

morgana% ping -c 2 remote1
64 bytes from 192.168.1.88: icmp_seq=0 ttl=64 time=1.8 ms
64 bytes from 192.168.1.88: icmp_seq=1 ttl=64 time=0.9 ms
2 packets transmitted, 2 packets received, 0% packet loss
```

Note that you won't be able to reach any other "host" belonging to the two networks because the packets are discarded by your computer after the address has been modified and the packet has been received. For example, a packet aimed at 192.168.0.32 will leave through `sn0` and reappear at `sn1` with a destination address of 192.168.1.32, which is not a local address for the host computer.

The Physical Transport of Packets

As far as data transport is concerned, the snull interfaces belong to the Ethernet class.

snull emulates Ethernet because the vast majority of existing networks —

– at least the segments that a workstation connects to -- are based on Ethernet technology, be it 10baseT, 100baseT, or gigabit. Additionally, the kernel offers some generalized support for Ethernet devices, and there's no reason not to use it. The advantage of being an Ethernet device is so strong that even the plip interface (the interface that uses the printer ports) declares itself as an Ethernet device.

The last advantage of using the Ethernet setup for snull is that you can run tcpdump on the interface to see the packets go by. Watching the interfaces with tcpdump can be a useful way to see how the two interfaces work. (Note that on 2.0 kernels, tcpdump will not work properly unless snull's interfaces show up as ethx. Load the driver with the eth=1 option to use the regular Ethernet names, rather than the default snx names.)

As was mentioned previously, snull only works with IP packets. This limitation is a result of the fact that snull snoops in the packets and even modifies them, in order for the code to work. The code modifies the source, destination, and checksum in the IP header of each packet without checking whether it actually conveys IP information. This quick-and-dirty data modification destroys non-IP packets. If you want to deliver other protocols through snull, you must modify the module's source code.

Connecting to the Kernel

We'll start looking at the structure of network drivers by dissecting the snull source. Keeping the source code for several drivers handy might help you follow the discussion and to see how real-world Linux network drivers operate. As a place to start, we suggest loopback.c, plip.c, and 3c509.c, in order of increasing complexity. Keeping skeleton.c handy might help as well, although this sample driver doesn't actually run. All these files live in drivers/net, within the kernel source tree.

Module Loading

When a driver module is loaded into a running kernel, it requests resources and offers facilities; there's nothing new in that. And there's also nothing new in the way resources are requested. The driver should probe for its device and its hardware location (I/O ports and IRQ line) -- but without registering them -- as described in "Installing an Interrupt Handler" in Chapter 9, "Interrupt Handling". The way a network driver is registered by its module initialization function is different from char and block drivers. Since there is no equivalent of major and minor numbers for network interfaces, a network driver does not request such a number. Instead, the driver inserts a data structure for each newly detected interface into a global list of network devices.

Each interface is described by a struct net_device item. The structures for sn0 and sn1, the two snullinterfaces, are declared like this:

```
struct net_device snull_devs[2] = {
```

```
        { init: snull_init, },   /* init, nothing more */
        { init: snull_init, }
    };
```

The initialization shown seems quite simple -- it sets only one field.
In fact, the net_device structure is huge, and we will be filling in other
pieces of it later on. But it is not helpful to cover the entire
structure at this point; instead, we will explain each field as it is
used. For the interested reader, the definition of the structure may be
found in <linux/netdevice.h>.

The first struct net_device field we will look at is name, which holds the
interface name (the string identifying the interface). The driver can
hardwire a name for the interface or it can allow dynamic assignment,
which works like this: if the name contains a %d format string, the
first available name found by replacing that string with a small integer
is used. Thus, eth%d is turned into the first available ethn name; the
first Ethernet interface is called eth0, and the others follow in
numeric order. The snullinterfaces are called sn0 and sn1 by default.
However, if eth=1 is specified at load time (causing the integer variable
snull_eth to be set to 1), snull_init uses dynamic assignment, as
follows:

```
    if (!snull_eth) { /* call them "sn0" and "sn1" */
        strcpy(snull_devs[0].name, "sn0");
        strcpy(snull_devs[1].name, "sn1");
    } else { /* use automatic assignment */
        strcpy(snull_devs[0].name, "eth%d");
        strcpy(snull_devs[1].name, "eth%d");
    }
```

The other field we initialized is init, a function pointer. Whenever you
register a device, the kernel asks the driver to initialize itself.
Initialization means probing for the physical interface and filling the
net_device structure with the proper values, as described in the following
section. If initialization fails, the structure is not linked to the
global list of network devices. This peculiar way of setting things up
is most useful during system boot; every driver tries to register its
own devices, but only devices that exist are linked to the list.

Because the real initialization is performed elsewhere, the
initialization function has little to do, and a single statement does
it:

```
    for (i=0; i<2;  i++)
        if ( (result = register_netdev(snull_devs + i)) )
            printk("snull: error %i registering device \"%s\"\n",
                    result, snull_devs[i].name);
        else device_present++;
```

Initializing Each Device

Probing for the device should be performed in the init function for the
interface (which is often called the "probe" function). The single
argument received by init is a pointer to the device being initialized;

its return value is either 0 or a negative error code, usually -ENODEV.

No real probing is performed for the snullinterface, because it is not bound to any hardware. When you write a real driver for a real interface, the usual rules for probing devices apply, depending on the peripheral bus you are using. Also, you should avoid registering I/O ports and interrupt lines at this point. Hardware registration should be delayed until device open time; this is particularly important if interrupt lines are shared with other devices. You don't want your interface to be called every time another device triggers an IRQ line just to reply "no, it's not mine."

The main role of the initialization routine is to fill in the dev structure for this device. Note that for network devices, this structure is always put together at runtime. Because of the way the network interface probing works, the dev structure cannot be set up at compile time in the same manner as a file_operations or block_device_operations structure. So, on exit from dev->init, the dev structure should be filled with correct values. Fortunately, the kernel takes care of some Ethernet-wide defaults through the function ether_setup, which fills several fields in struct net_device.

The core of snull_init is as follows:

```
ether_setup(dev); /* assign some of the fields */

dev->open            = snull_open;
dev->stop            = snull_release;
dev->set_config      = snull_config;
dev->hard_start_xmit = snull_tx;
dev->do_ioctl        = snull_ioctl;
dev->get_stats       = snull_stats;
dev->rebuild_header  = snull_rebuild_header;
dev->hard_header     = snull_header;
#ifdef HAVE_TX_TIMEOUT
dev->tx_timeout      = snull_tx_timeout;
dev->watchdog_timeo = timeout;
#endif
/* keep the default flags, just add NOARP */
dev->flags           |= IFF_NOARP;
dev->hard_header_cache = NULL;      /* Disable caching */
SET_MODULE_OWNER(dev);
```

The single unusual feature of the code is setting IFF_NOARP in the flags. This specifies that the interface cannot use ARP, the Address Resolution Protocol. ARP is a low-level Ethernet protocol; its job is to turn IP addresses into Ethernet Medium Access Control (MAC) addresses. Since the "remote" systems simulated by snull do not really exist, there is nobody available to answer ARP requests for them. Rather than complicate snull with the addition of an ARP implementation, we chose to mark the interface as being unable to handle that protocol. The assignment to hard_header_cache is there for a similar reason: it disables the caching of the (nonexistent) ARP replies on this interface. This topic is discussed in detail later in this chapter in "MAC Address Resolution".

The initialization code also sets a couple of fields (tx_timeout and watchdog_timeo) that relate to the handling of transmission timeouts. We

will cover this topic thoroughly later in this chapter in "Transmission Timeouts".

Finally, this code calls SET_MODULE_OWNER, which initializes the owner field of the net_device structure with a pointer to the module itself. The kernel uses this information in exactly the same way it uses the owner field of the file_operations structure -- to maintain the module's usage count.

We'll look now at one more struct net_device field, priv. Its role is similar to that of the private_data pointer that we used for char drivers. Unlike fops->private_data, this priv pointer is allocated at initialization time instead of open time, because the data item pointed to by priv usually includes the statistical information about interface activity. It's important that statistical information always be available, even when the interface is down, because users may want to display the statistics at any time by calling ifconfig. The memory wasted by allocating priv during initialization instead of on open is irrelevant because most probed interfaces are constantly up and running in the system. The snull module declares a snull_priv data structure to be used for priv:

```
struct snull_priv {
    struct net_device_stats stats;
    int status;
    int rx_packetlen;
    u8 *rx_packetdata;
    int tx_packetlen;
    u8 *tx_packetdata;
    struct sk_buff *skb;
    spinlock_t lock;
};
```

The structure includes an instance of struct net_device_stats, which is the standard place to hold interface statistics. The following lines in snull_init allocate and initialize dev->priv:

```
dev->priv = kmalloc(sizeof(struct snull_priv), GFP_KERNEL);
if (dev->priv == NULL)
    return -ENOMEM;
memset(dev->priv, 0, sizeof(struct snull_priv));
spin_lock_init(& ((struct snull_priv *) dev->priv)->lock);
```

Module Unloading

Nothing special happens when the module is unloaded. The module cleanup function simply unregisters the interfaces from the list after releasing memory associated with the private structure:

```
void snull_cleanup(void)
{
    int i;

    for (i=0; i<2; i++) {
        kfree(snull_devs[i].priv);
```

```
        unregister_netdev(snull_devs + i);
    }
    return;
}
```

Modularized and Nonmodularized Drivers

Although char and block drivers are the same regardless of whether they're modular or linked into the kernel, that's not the case for network drivers.

When a driver is linked directly into the Linux kernel, it doesn't declare its own `net_device` structures; the structures declared in drivers/net/Space.c are used instead. Space.c declares a linked list of all the network devices, both driver-specific structures like `plip1` and general-purpose `eth` devices. Ethernet drivers don't care about their `net_device` structures at all, because they use the general-purpose structures. Such general `eth` device structures declare ethif_probe as their init function. A programmer inserting a new Ethernet interface in the mainstream kernel needs only to add a call to the driver's initialization function to ethif_probe. Authors of non-`eth` drivers, on the other hand, insert their `net_device` structures in Space.c. In both cases only the source file Space.c has to be modified if the driver must be linked to the kernel proper.

At system boot, the network initialization code loops through all the `net_device` structures and calls their probing (`dev->init`) functions by passing them a pointer to the device itself. If the probe function succeeds, the kernel initializes the next available `net_device` structure to use that interface. This way of setting up drivers permits incremental assignment of devices to the names `eth0`, `eth1`, and so on, without changing the `name` field of each device.

When a modularized driver is loaded, on the other hand, it declares its own `net_device` structures (as we have seen in this chapter), even if the interface it controls is an Ethernet interface.

The curious reader can learn more about interface initialization by looking at Space.c and net_init.c.

The net_device Structure in Detail

The `net_device` structure is at the very core of the network driver layer and deserves a complete description. At a first reading, however, you can skip this section, because you don't need a thorough understanding of the structure to get started. This list describes all the fields, but more to provide a reference than to be memorized. The rest of this chapter briefly describes each field as soon as it is used in the sample code, so you don't need to keep referring back to this section.

`struct net_device` can be conceptually divided into two parts: visible and invisible. The visible part of the structure is made up of the fields that can be explicitly assigned in static `net_device` structures. All structures in drivers/net/Space.c are initialized in this way, without

using the tagged syntax for structure initialization. The remaining fields are used internally by the network code and usually are not initialized at compilation time, not even by tagged initialization. Some of the fields are accessed by drivers (for example, the ones that are assigned at initialization time), while some shouldn't be touched.

The Visible Head

The first part of `struct net_device` is composed of the following fields, in this order:

`char name[IFNAMSIZ];`

> The name of the device. If the name contains a `%d` format string, the first available device name with the given base is used; assigned numbers start at zero.

`unsigned long rmem_end;`
`unsigned long rmem_start;`
`unsigned long mem_end;`
`unsigned long mem_start;`

> Device memory information. These fields hold the beginning and ending addresses of the shared memory used by the device. If the device has different receive and transmit memories, the `mem` fields are used for transmit memory and the `rmem` fields for receive memory. `mem_start` and `mem_end` can be specified on the kernel command line at system boot, and their values are retrieved by ifconfig. The `rmem` fields are never referenced outside of the driver itself. By convention, the `end` fields are set so that `end - start` is the amount of available on-board memory.

`unsigned long base_addr;`

> The I/O base address of the network interface. This field, like the previous ones, is assigned during device probe. The ifconfig command can be used to display or modify the current value. The `base_addr` can be explicitly assigned on the kernel command line at system boot or at load time. The field is not used by the kernel, like the memory fields shown previously.

`unsigned char irq;`

> The assigned interrupt number. The value of `dev->irq` is printed by ifconfig when interfaces are listed. This value can usually be set at boot or load time and modified later using ifconfig.

`unsigned char if_port;`

> Which port is in use on multiport devices. This field is used, for example, with devices that support both coaxial (`IF_PORT_10BASE2`) and twisted-pair (`IF_PORT_10BASET`) Ethernet connections. The full set of known port types is defined in `<linux/netdevice.h>`.

`unsigned char dma;`

The DMA channel allocated by the device. The field makes sense only with some peripheral buses, like ISA. It is not used outside of the device driver itself, but for informational purposes (in ifconfig).

unsigned long state;

Device state. The field includes several flags. Drivers do not normally manipulate these flags directly; instead, a set of utility functions has been provided. These functions will be discussed shortly when we get into driver operations.

struct net_device *next;

Pointer to the next device in the global linked list. This field shouldn't be touched by the driver.

int (*init)(struct net_device *dev);

The initialization function, described earlier.

The Hidden Fields

The net_device structure includes many additional fields, which are usually assigned at device initialization. Some of these fields convey information about the interface, while some exist only for the benefit of the driver (i.e., they are not used by the kernel); other fields, most notably the device methods, are part of the kernel-driver interface.

We will list the three groups separately, independent of the actual order of the fields, which is not significant.

Interface information

Most of the information about the interface is correctly set up by the function ether_setup. Ethernet cards can rely on this general-purpose function for most of these fields, but the flags and dev_addr fields are device specific and must be explicitly assigned at initialization time.

Some non-Ethernet interfaces can use helper functions similar to ether_setup. drivers/net/net_init.c exports a number of such functions, including the following:

void ltalk_setup(struct net_device *dev);

Sets up the fields for a LocalTalk device.

void fc_setup(struct net_device *dev);

Initializes for fiber channel devices.

void fddi_setup(struct net_device *dev);

Configures an interface for a Fiber Distributed Data Interface (FDDI) network.

```
void hippi_setup(struct net_device *dev);
```

> Prepares fields for a High-Performance Parallel Interface (HIPPI)
> high-speed interconnect driver.

```
void tr_configure(struct net_device *dev);
```

> Handles setup for token ring network interfaces. Note that the 2.4
> kernel also exports a function tr_setup, which, interestingly, does
> nothing at all.

Most devices will be covered by one of these classes. If yours is
something radically new and different, however, you will need to assign
the following fields by hand.

```
unsigned short hard_header_len;
```

> The hardware header length, that is, the number of octets that lead
> the transmitted packet before the IP header, or other protocol
> information. The value of hard_header_len is 14 (ETH_HLEN) for Ethernet
> interfaces.

```
unsigned mtu;
```

> The maximum transfer unit (MTU). This field is used by the network
> layer to drive packet transmission. Ethernet has an MTU of 1500
> octets (ETH_DATA_LEN).

```
unsigned long tx_queue_len;
```

> The maximum number of frames that can be queued on the device's
> transmission queue. This value is set to 100 by ether_setup, but
> you can change it. For example, plip uses 10 to avoid wasting
> system memory (plip has a lower throughput than a real Ethernet
> interface).

```
unsigned short type;
```

> The hardware type of the interface. The type field is used by ARP
> to determine what kind of hardware address the interface supports.
> The proper value for Ethernet interfaces is ARPHRD_ETHER, and that is
> the value set by ether_setup. The recognized types are defined in
> <linux/if_arp.h>.

```
unsigned char addr_len;
unsigned char broadcast[MAX_ADDR_LEN];
unsigned char dev_addr[MAX_ADDR_LEN];
```

> Hardware (MAC) address length and device hardware addresses. The
> Ethernet address length is six octets (we are referring to the
> hardware ID of the interface board), and the broadcast address is
> made up of six 0xff octets; ether_setup arranges for these values
> to be correct. The device address, on the other hand, must be read
> from the interface board in a device-specific way, and the driver
> should copy it to dev_addr. The hardware address is used to generate
> correct Ethernet headers before the packet is handed over to the
> driver for transmission. The snulldevice doesn't use a physical

interface, and it invents its own hardware address.

`unsigned short flags;`

Interface flags, detailed next.

The `flags` field is a bit mask including the following bit values. The `IFF_` prefix stands for "interface flags." Some flags are managed by the kernel, and some are set by the interface at initialization time to assert various capabilities and other features of the interface. The valid flags, which are defined in `<linux/if.h>`, are as follows:

IFF_UP

This flag is read-only for the driver. The kernel turns it on when the interface is active and ready to transfer packets.

IFF_BROADCAST

This flag states that the interface allows broadcasting. Ethernet boards do.

IFF_DEBUG

This marks debug mode. The flag can be used to control the verbosity of your printk calls or for other debugging purposes. Although no official driver currently uses this flag, it can be set and reset by user programs via ioctl, and your driver can use it. The misc-progs/netifdebug program can be used to turn the flag on and off.

IFF_LOOPBACK

This flag should be set only in the loopback interface. The kernel checks for `IFF_LOOPBACK` instead of hardwiring the `lo` name as a special interface.

IFF_POINTOPOINT

This flag signals that the interface is connected to a point-to-point link. It is set by ifconfig. For example, plip and the PPP driver have it set.

IFF_NOARP

This means that the interface can't perform ARP. For example, point-to-point interfaces don't need to run ARP, which would only impose additional traffic without retrieving useful information. snull runs without ARP capabilities, so it sets the flag.

IFF_PROMISC

This flag is set to activate promiscuous operation. By default, Ethernet interfaces use a hardware filter to ensure that they receive broadcast packets and packets directed to that interface's hardware address only. Packet sniffers such as tcpdump set promiscuous mode on the interface in order to retrieve all packets

that travel on the interface's transmission medium.

IFF_MULTICAST

This flag is set by interfaces that are capable of multicast transmission. ether_setup sets IFF_MULTICAST by default, so if your driver does not support multicast, it must clear the flag at initialization time.

IFF_ALLMULTI

This flag tells the interface to receive all multicast packets. The kernel sets it when the host performs multicast routing, only if IFF_MULTICAST is set. IFF_ALLMULTI is read-only for the interface. We'll see the multicast flags used in "Multicasting" later in this chapter.

IFF_MASTER
IFF_SLAVE

These flags are used by the load equalization code. The interface driver doesn't need to know about them.

IFF_PORTSEL
IFF_AUTOMEDIA

These flags signal that the device is capable of switching between multiple media types, for example, unshielded twisted pair (UTP) versus coaxial Ethernet cables. If IFF_AUTOMEDIA is set, the device selects the proper medium automatically.

IFF_DYNAMIC

This flag indicates that the address of this interface can change; used with dialup devices.

IFF_RUNNING

This flag indicates that the interface is up and running. It is mostly present for BSD compatibility; the kernel makes little use of it. Most network drivers need not worry about IFF_RUNNING.

IFF_NOTRAILERS

This flag is unused in Linux, but it exists for BSD compatibility.

When a program changes IFF_UP, the open or stop device method is called. When IFF_UP or any other flag is modified, the set_multicast_list method is invoked. If the driver needs to perform some action because of a modification in the flags, it must take that action in set_multicast_list. For example, when IFF_PROMISC is set or reset, set_multicast_list must notify the onboard hardware filter. The responsibilities of this device method are outlined in "Multicasting".

The device methods

As happens with the char and block drivers, each network device declares

the functions that act on it. Operations that can be performed on
network interfaces are listed in this section. Some of the operations
can be left NULL, and some are usually untouched because ether_setup
assigns suitable methods to them.

Device methods for a network interface can be divided into two groups:
fundamental and optional. Fundamental methods include those that are
needed to be able to use the interface; optional methods implement more
advanced functionalities that are not strictly required. The following
are the fundamental methods:

int (*open)(struct net_device *dev);

> Opens the interface. The interface is opened whenever ifconfig
> activates it. The open method should register any system resource
> it needs (I/O ports, IRQ, DMA, etc.), turn on the hardware, and
> increment the module usage count.

int (*stop)(struct net_device *dev);

> Stops the interface. The interface is stopped when it is brought
> down; operations performed at open time should be reversed.

int (*hard_start_xmit) (struct sk_buff *skb, struct net_device *dev);

> This method initiates the transmission of a packet. The full packet
> (protocol headers and all) is contained in a socket buffer (sk_buff)
> structure. Socket buffers are introduced later in this chapter.

int (*hard_header) (struct sk_buff *skb, struct net_device *dev, unsigned short type,
void *daddr, void *saddr, unsigned len);

> This function builds the hardware header from the source and
> destination hardware addresses that were previously retrieved; its
> job is to organize the information passed to it as arguments into
> an appropriate, device-specific hardware header. eth_header is the
> default function for Ethernet-like interfaces, and ether_setup
> assigns this field accordingly.

int (*rebuild_header)(struct sk_buff *skb);

> This function is used to rebuild the hardware header before a
> packet is transmitted. The default function used by Ethernet
> devices uses ARP to fill the packet with missing information. The
> rebuild_header method is used rarely in the 2.4 kernel; hard_header
> is used instead.

void (*tx_timeout)(struct net_device *dev);

> This method is called when a packet transmission fails to complete
> within a reasonable period, on the assumption that an interrupt has
> been missed or the interface has locked up. It should handle the
> problem and resume packet transmission.

struct net_device_stats *(*get_stats)(struct net_device *dev);

> Whenever an application needs to get statistics for the interface,

this method is called. This happens, for example, when ifconfig or netstat -i is run. A sample implementation for snull is introduced in "Statistical Information" later in this chapter.

int (*set_config)(struct net_device *dev, struct ifmap *map);

Changes the interface configuration. This method is the entry point for configuring the driver. The I/O address for the device and its interrupt number can be changed at runtime using set_config. This capability can be used by the system administrator if the interface cannot be probed for. Drivers for modern hardware normally do not need to implement this method.

The remaining device operations may be considered optional.

int (*do_ioctl)(struct net_device *dev, struct ifreq *ifr, int cmd);

Perform interface-specific ioctl commands. Implementation of those commands is described later in "Custom ioctl Commands". The corresponding field in struct net_device can be left as NULL if the interface doesn't need any interface-specific commands.

void (*set_multicast_list)(struct net_device *dev);

This method is called when the multicast list for the device changes and when the flags change. See "Multicasting" for further details and a sample implementation.

int (*set_mac_address)(struct net_device *dev, void *addr);

This function can be implemented if the interface supports the ability to change its hardware address. Many interfaces don't support this ability at all. Others use the default eth_mac_addr implementation (from drivers/net/net_init.c). eth_mac_addr only copies the new address into dev->dev_addr, and it will only do so if the interface is not running. Drivers that use eth_mac_addr should set the hardware MAC address from dev->dev_addr when they are configured.

int (*change_mtu)(struct net_device *dev, int new_mtu);

This function is in charge of taking action if there is a change in the MTU (maximum transfer unit) for the interface. If the driver needs to do anything particular when the MTU is changed, it should declare its own function; otherwise, the default will do the right thing. snull has a template for the function if you are interested.

int (*header_cache) (struct neighbour *neigh, struct hh_cache *hh);

header_cache is called to fill in the hh_cache structure with the results of an ARP query. Almost all drivers can use the default eth_header_cache implementation.

int (*header_cache_update) (struct hh_cache *hh, struct net_device *dev, unsigned char *haddr);

This method updates the destination address in the hh_cache

> structure in response to a change. Ethernet devices use
> eth_header_cache_update.

int (*hard_header_parse) (struct sk_buff *skb, unsigned char *haddr);

> The hard_header_parse method extracts the source address from the
> packet contained in skb, copying it into the buffer at haddr. The
> return value from the function is the length of that address.
> Ethernet devices normally use eth_header_parse.

Utility fields

The remaining struct net_device data fields are used by the interface to
hold useful status information. Some of the fields are used by ifconfig
and netstat to provide the user with information about the current
configuration. An interface should thus assign values to these fields.

unsigned long trans_start;
unsigned long last_rx;

> Both of these fields are meant to hold a jiffies value. The driver
> is responsible for updating these values when transmission begins
> and when a packet is received, respectively. The trans_start value is
> used by the networking subsystem to detect transmitter lockups.
> last_rx is currently unused, but the driver should maintain this
> field anyway to be prepared for future use.

int watchdog_timeo;

> The minimum time (in jiffies) that should pass before the
> networking layer decides that a transmission timeout has occurred
> and calls the driver's tx_timeout function.

void *priv;

> The equivalent of filp->private_data. The driver owns this pointer and
> can use it at will. Usually the private data structure includes a
> struct net_device_stats item. The field is used in "Initializing Each
> Device", later in this chapter.

struct dev_mc_list *mc_list;
int mc_count;

> These two fields are used in handling multicast transmission.
> mc_count is the count of items in mc_list. See "Multicasting" for
> further details.

spinlock_t xmit_lock;
int xmit_lock_owner;

> The xmit_lock is used to avoid multiple simultaneous calls to the
> driver's hard_start_xmit function. xmit_lock_owner is the number of
> the CPU that has obtained xmit_lock. The driver should make no
> changes to these fields.

struct module *owner;

The module that "owns" this device structure; it is used to maintain the use count for the module.

There are other fields in `struct net_device`, but they are not used by network drivers.

Opening and Closing

Our driver can probe for the interface at module load time or at kernel boot. Before the interface can carry packets, however, the kernel must open it and assign an address to it. The kernel will open or close an interface in response to the ifconfig command.

When ifconfig is used to assign an address to the interface, it performs two tasks. First, it assigns the address by means of `ioctl(SIOCSIFADDR)` (Socket I/O Control Set Interface Address). Then it sets the `IFF_UP` bit in `dev->flag` by means of `ioctl(SIOCSIFFLAGS)` (Socket I/O Control Set Interface Flags) to turn the interface on.

As far as the device is concerned, `ioctl(SIOCSIFADDR)` does nothing. No driver function is invoked -- the task is device independent, and the kernel performs it. The latter command (`ioctl(SIOCSIFFLAGS)`), though, calls the open method for the device.

Similarly, when the interface is shut down, ifconfig uses `ioctl(SIOCSIFFLAGS)` to clear `IFF_UP`, and the stop method is called.

Both device methods return 0 in case of success and the usual negative value in case of error.

As far as the actual code is concerned, the driver has to perform many of the same tasks as the char and block drivers do. open requests any system resources it needs and tells the interface to come up; stop shuts down the interface and releases system resources. There are a couple of additional steps to be performed, however.

First, the hardware address needs to be copied from the hardware device to `dev->dev_addr` before the interface can communicate with the outside world. The hardware address can be assigned at probe time or at open time, at the driver's will. The snull software interface assigns it from within open; it just fakes a hardware number using an ASCII string of length `ETH_ALEN`, the length of Ethernet hardware addresses.

The open method should also start the interface's transmit queue (allow it to accept packets for transmission) once it is ready to start sending data. The kernel provides a function to start the queue:

```
void netif_start_queue(struct net_device *dev);
```

The open code for snull looks like the following:

```
int snull_open(struct net_device *dev)
{
    MOD_INC_USE_COUNT;
```

```
        /* request_region(), request_irq(), .... (like fops->open) */

        /*
         * Assign the hardware address of the board: use "\0SNULx", where
         * x is 0 or 1. The first byte is '\0' to avoid being a multicast
         * address (the first byte of multicast addrs is odd).
         */
        memcpy(dev->dev_addr, "\0SNUL0", ETH_ALEN);
        dev->dev_addr[ETH_ALEN-1] += (dev - snull_devs); /* the number */

        netif_start_queue(dev);
        return 0;
    }
```

As you can see, in the absence of real hardware, there is little to do
in the open method. The same is true of the stop method; it just
reverses the operations of open. For this reason the function
implementing stop is often called close or release.

```
    int snull_release(struct net_device *dev)
    {
        /* release ports, irq and such -- like fops->close */

        netif_stop_queue(dev); /* can't transmit any more */
        MOD_DEC_USE_COUNT;
        return 0;
    }
```

The function:

```
    void netif_stop_queue(struct net_device *dev);
```

is the opposite of netif_start_queue; it marks the device as being
unable to transmit any more packets. The function must be called when
the interface is closed (in the stop method) but can also be used to
temporarily stop transmission, as explained in the next section.

Packet Transmission

The most important tasks performed by network interfaces are data
transmission and reception. We'll start with transmission because it is
slightly easier to understand.

Whenever the kernel needs to transmit a data packet, it calls the
hard_start_transmit method to put the data on an outgoing queue. Each
packet handled by the kernel is contained in a socket buffer structure
(struct sk_buff), whose definition is found in <linux/skbuff.h>. The
structure gets its name from the Unix abstraction used to represent a
network connection, the socket. Even if the interface has nothing to do
with sockets, each network packet belongs to a socket in the higher
network layers, and the input/output buffers of any socket are lists of
struct sk_buff structures. The same sk_buff structure is used to host
network data throughout all the Linux network subsystems, but a socket
buffer is just a packet as far as the interface is concerned.

A pointer to `sk_buff` is usually called `skb`, and we follow this practice both in the sample code and in the text.

The socket buffer is a complex structure, and the kernel offers a number of functions to act on it. The functions are described later in "The Socket Buffers"; for now a few basic facts about `sk_buff` are enough for us to write a working driver.

The socket buffer passed to hard_start_xmitcontains the physical packet as it should appear on the media, complete with the transmission-level headers. The interface doesn't need to modify the data being transmitted. `skb->data` points to the packet being transmitted, and `skb->len` is its length, in octets.

The snull packet transmission code is follows; the physical transmission machinery has been isolated in another function because every interface driver must implement it according to the specific hardware being driven.

```
int snull_tx(struct sk_buff *skb, struct net_device *dev)
{
    int len;
    char *data;
    struct snull_priv *priv = (struct snull_priv *) dev->priv;
    len = skb->len < ETH_ZLEN ? ETH_ZLEN : skb->len;
    data = skb->data;
    dev->trans_start = jiffies; /* save the timestamp */

    /* Remember the skb, so we can free it at interrupt time */
    priv->skb = skb;

    /* actual delivery of data is device specific, and not shown here */
    snull_hw_tx(data, len, dev);

    return 0; /* Our simple device cannot fail */
}
```

The transmission function thus performs only some sanity checks on the packet and transmits the data through the hardware-related function. That function (snull_hw_tx) is omitted here since it is entirely occupied with implementing the trickery of the snull device (including manipulating the source and destination addresses) and has little of interest to authors of real network drivers. It is present, of course, in the sample source for those who want to go in and see how it works.

Controlling Transmission Concurrency

The hard_start_xmit function is protected from concurrent calls by a spinlock (`xmit_lock`) in the `net_device` structure. As soon as the function returns, however, it may be called again. The function returns when the software is done instructing the hardware about packet transmission, but hardware transmission will likely not have been completed. This is not an issue with snull, which does all of its work using the CPU, so packet transmission is complete before the transmission function returns.

Real hardware interfaces, on the other hand, transmit packets

asynchronously and have a limited amount of memory available to store outgoing packets. When that memory is exhausted (which, for some hardware, will happen with a single outstanding packet to transmit), the driver will need to tell the networking system not to start any more transmissions until the hardware is ready to accept new data.

This notification is accomplished by calling netif_stop_queue, the function introduced earlier to stop the queue. Once your driver has stopped its queue, it must arrange to restart the queue at some point in the future, when it is again able to accept packets for transmission. To do so, it should call:

```
void netif_wake_queue(struct net_device *dev);
```

This function is just like netif_start_queue, except that it also pokes the networking system to make it start transmitting packets again.

Most modern network interfaces maintain an internal queue with multiple packets to transmit; in this way they can get the best performance from the network. Network drivers for these devices support having multiple transmisions outstanding at any given time, but device memory can fill up whether or not the hardware supports multiple outstanding transmission. Whenever device memory fills to the point that there is no room for the largest possible packet, the driver should stop the queue until space becomes available again.

Transmission Timeouts

Most drivers that deal with real hardware have to be prepared for that hardware to fail to respond occasionally. Interfaces can forget what they are doing, or the system can lose an interrupt. This sort of problem is common with some devices designed to run on personal computers.

Many drivers handle this problem by setting timers; if the operation has not completed by the time the timer expires, something is wrong. The network system, as it happens, is essentially a complicated assembly of state machines controlled by a mass of timers. As such, the networking code is in a good position to detect transmission timeouts automatically.

Thus, network drivers need not worry about detecting such problems themselves. Instead, they need only set a timeout period, which goes in the watchdog_timeo field of the net_device structure. This period, which is in jiffies, should be long enough to account for normal transmission delays (such as collisions caused by congestion on the network media).

If the current system time exceeds the device's trans_start time by at least the timeout period, the networking layer will eventually call the driver's tx_timeout method. That method's job is to do whatever is needed to clear up the problem and to ensure the proper completion of any transmissions that were already in progress. It is important, in particular, that the driver not lose track of any socket buffers that have been entrusted to it by the networking code.

snull has the ability to simulate transmitter lockups, which is controlled by two load-time parameters:

```
static int lockup = 0;
MODULE_PARM(lockup, "i");

#ifdef HAVE_TX_TIMEOUT
static int timeout = SNULL_TIMEOUT;
MODULE_PARM(timeout, "i");
#endif
```

If the driver is loaded with the parameter lockup=n, a lockup will be simulated once every n packets transmitted, and the watchdog_timeo field will be set to the given timeout value. When simulating lockups, snull also calls netif_stop_queue to prevent other transmission attempts from occurring.

The snull transmission timeout handler looks like this:

```
void snull_tx_timeout (struct net_device *dev)
{
    struct snull_priv *priv = (struct snull_priv *) dev->priv;

    PDEBUG("Transmit timeout at %ld, latency %ld\n", jiffies,
                    jiffies - dev->trans_start);
    priv->status = SNULL_TX_INTR;
    snull_interrupt(0, dev, NULL);
    priv->stats.tx_errors++;
    netif_wake_queue(dev);
    return;
}
```

When a transmission timeout happens, the driver must mark the error in the interface statistics and arrange for the device to be reset to a sane state so that new packets can be transmitted. When a timeout happens in snull, the driver calls snull_interrupt to fill in the "missing" interrupt and restarts the transmit queue with netif_wake_queue.

Packet Reception

Receiving data from the network is trickier than transmitting it because an sk_buff must be allocated and handed off to the upper layers from within an interrupt handler. The usual way to receive a packet is through an interrupt, unless the interface is a purely software one like snull or the loopback interface. Although it is possible to write polling drivers, and a few exist in the official kernel, interrupt-driven operation is much better, both in terms of data throughput and computational demands. Because most network interfaces are interrupt driven, we won't talk about the polling implementation, which just exploits kernel timers.

The implementation of snull separates the "hardware" details from the device-independent housekeeping. The function snull_rx is thus called after the hardware has received the packet and it is already in the

computer's memory. snull_rx receives a pointer to the data and the
length of the packet; its sole responsibility is to send the packet and
some additional information to the upper layers of networking code. This
code is independent of the way the data pointer and length are obtained.

```c
void snull_rx(struct net_device *dev, int len, unsigned char *buf)
{
    struct sk_buff *skb;
    struct snull_priv *priv = (struct snull_priv *) dev->priv;

    /*
     * The packet has been retrieved from the transmission
     * medium. Build an skb around it, so upper layers can handle it
     */
    skb = dev_alloc_skb(len+2);
    if (!skb) {
        printk("snull rx: low on mem - packet dropped\n");
        priv->stats.rx_dropped++;
        return;
    }
    memcpy(skb_put(skb, len), buf, len);

    /* Write metadata, and then pass to the receive level */
    skb->dev = dev;
    skb->protocol = eth_type_trans(skb, dev);
    skb->ip_summed = CHECKSUM_UNNECESSARY; /* don't check it */
    priv->stats.rx_packets++;
    priv->stats.rx_bytes += len;
    netif_rx(skb);
    return;
}
```

The function is sufficiently general to act as a template for any
network driver, but some explanation is necessary before you can reuse
this code fragment with confidence.

The first step is to allocate a buffer to hold the packet. Note that the
buffer allocation function (dev_alloc_skb) needs to know the data
length. The information is used by the function to allocate space for
the buffer. dev_alloc_skbcalls kmalloc with atomic priority; it can thus
be used safely at interrupt time. The kernel offers other interfaces to
socket-buffer allocation, but they are not worth introducing here;
socket buffers are explained in detail in "The Socket Buffers", later in
this chapter.

Once there is a valid skb pointer, the packet data is copied into the
buffer by calling memcpy; the skb_put function updates the end-of-data
pointer in the buffer and returns a pointer to the newly created space.

If you are writing a high-performance driver for an interface that can
do full bus-mastering I/O, there is a possible optimization that is
worth considering here. Some drivers allocate socket buffers for
incoming packets prior to their reception, then instruct the interface
to place the packet data directly into the socket buffer's space. The
networking layer cooperates with this strategy by allocating all socket
buffers in DMA-capable space. Doing things this way avoids the need for
a separate copy operation to fill the socket buffer, but requires being
careful with buffer sizes because you won't know in advance how big the

incoming packet is. The implementation of a change_mtu method is also important in this situation, since it allows the driver to respond to a change in the maximum packet size.

The network layer needs to have some information spelled out before it will be able to make sense of the packet. To this end, the `dev` and `protocol` fields must be assigned before the buffer is passed upstairs. Then we need to specify how checksumming is to be performed or has been performed on the packet (snull does not need to perform any checksums). The possible policies for `skb->ip_summed` are as follows:

CHECKSUM_HW

> The device has already performed checksums in hardware. An example of a hardware checksum is the SPARC HME interface.

CHECKSUM_NONE

> Checksums are still to be verified, and the task must be accomplished by system software. This is the default in newly allocated buffers.

CHECKSUM_UNNECESSARY

> Don't do any checksums. This is the policy in snull and in the loopback interface.

Finally, the driver updates its statistics counter to record that a packet has been received. The statistics structure is made up of several fields; the most important are `rx_packets`, `rx_bytes`, `tx_packets`, and `tx_bytes`, which contain the number of packets received and transmitted and the total number of octets transferred. All the fields are thoroughly described in "Statistical Information" later in this chpater.

The last step in packet reception is performed by netif_rx, which hands off the socket buffer to the upper layers.

The Interrupt Handler

Most hardware interfaces are controlled by means of an interrupt handler. The interface interrupts the processor to signal one of two possible events: a new packet has arrived or transmission of an outgoing packet is complete. This generalization doesn't always apply, but it does account for all the problems related to asynchronous packet transmission. Parallel Line Internet Protocol (PLIP) and Point-to-Point Protocol (PPP) are examples of interfaces that don't fit this generalization. They deal with the same events, but the low-level interrupt handling is slightly different.

The usual interrupt routine can tell the difference between a new-packet-arrived interrupt and a done-transmitting notification by checking a status register found on the physical device. The snull interface works similarly, but its status word is implemented in software and lives in `dev->priv`. The interrupt handler for a network interface looks like this:

```
void snull_interrupt(int irq, void *dev_id, struct pt_regs *regs)
{
    int statusword;
    struct snull_priv *priv;
    /*
     * As usual, check the "device" pointer for shared handlers.
     * Then assign "struct device *dev"
     */
    struct net_device *dev = (struct net_device *)dev_id;
    /* ... and check with hw if it's really ours */

    if (!dev /*paranoid*/ ) return;

    /* Lock the device */
    priv = (struct snull_priv *) dev->priv;
    spin_lock(&priv->lock);

    /* retrieve statusword: real netdevices use I/O instructions */
    statusword = priv->status;
    if (statusword & SNULL_RX_INTR) {
        /* send it to snull_rx for handling */
        snull_rx(dev, priv->rx_packetlen, priv->rx_packetdata);
    }
    if (statusword & SNULL_TX_INTR) {
        /* a transmission is over: free the skb */
        priv->stats.tx_packets++;
        priv->stats.tx_bytes += priv->tx_packetlen;
        dev_kfree_skb(priv->skb);
    }

    /* Unlock the device and we are done */
    spin_unlock(&priv->lock);
    return;
}
```

The handler's first task is to retrieve a pointer to the correct struct
net_device. This pointer usually comes from the dev_id pointer received as
an argument.

The interesting part of this handler deals with the "transmission done"
situation. In this case, the statistics are updated, and dev_kfree_skb
is called to return the (no longer needed) socket buffer to the system.
If your driver has temporarily stopped the transmission queue, this is
the place to restart it with netif_wake_queue.

Packet reception, on the other hand, doesn't need any special interrupt
handling. Calling snull_rx (which we have already seen) is all that's
required.

Changes in Link State

Network connections, by definition, deal with the world outside the
local system. They are thus often affected by outside events, and they
can be transient things. The networking subsystem needs to know when
network links go up or down, and it provides a few functions that the
driver may use to convey that information.

Most networking technologies involving an actual, physical connection
provide a carrier state; the presence of the carrier means that the

hardware is present and ready to function. Ethernet adapters, for example, sense the carrier signal on the wire; when a user trips over the cable, that carrier vanishes, and the link goes down. By default, network devices are assumed to have a carrier signal present. The driver can change that state explicitly, however, with these functions:

```
void netif_carrier_off(struct net_device *dev);
void netif_carrier_on(struct net_device *dev);
```

If your driver detects a lack of carrier on one of its devices, it should call netif_carrier_off to inform the kernel of this change. When the carrier returns, netif_carrier_on should be called. Some drivers also call netif_carrier_off when making major configuration changes (such as media type); once the adapter has finished resetting itself, the new carrier will be detected and traffic can resume.

An integer function also exsists:

```
int netif_carrier_ok(struct net_device *dev);
```

This can be used to test the current carrier state (as reflected in the device structure).

The Socket Buffers

We've now discussed most of the issues related to network interfaces. What's still missing is some more detailed discussion of the sk_buff structure. The structure is at the core of the network subsystem of the Linux kernel, and we now introduce both the main fields of the structure and the functions used to act on it.

Although there is no strict need to understand the internals of sk_buff, the ability to look at its contents can be helpful when you are tracking down problems and when you are trying to optimize the code. For example, if you look in loopback.c, you'll find an optimization based on knowledge of the sk_buff internals. The usual warning applies here: if you write code that takes advantage of knowledge of the sk_buff structure, you should be prepared to see it break with future kernel releases. Still, sometimes the performance advantages justify the additional maintenance cost.

We are not going to describe the whole structure here, just the fields that might be used from within a driver. If you want to see more, you can look at <linux/skbuff.h>, where the structure is defined and the functions are prototyped. Additional details about how the fields and functions are used can be easily retrieved by grepping in the kernel sources.

The Important Fields

The fields introduced here are the ones a driver might need to access. They are listed in no particular order.

```
struct net_device *rx_dev;
struct net_device *dev;
```

The devices receiving and sending this buffer, respectively.

```
union { /* ... */ } h;
union { /* ... */ } nh;
union { /*... */} mac;
```

>   Pointers to the various levels of headers contained within the
>   packet. Each field of the unions is a pointer to a different type
>   of data structure. `h` hosts pointers to transport layer headers (for
>   example, `struct tcphdr *th`); `nh` includes network layer headers (such
>   as `struct iphdr *iph`); and `mac` collects pointers to link layer headers
>   (such as `struct ethdr *ethernet`).
>
>   If your driver needs to look at the source and destination
>   addresses of a TCP packet, it can find them in `skb->h.th`. See the
>   header file for the full set of header types that can be accessed
>   in this way.
>
>   Note that network drivers are responsible for setting the `mac`
>   pointer for incoming packets. This task is normally handled by
>   ether_type_trans, but non-Ethernet drivers will have to set `skb->mac.raw` directly, as shown later in "Non-Ethernet Headers".

```
unsigned char *head;
unsigned char *data;
unsigned char *tail;
unsigned char *end;
```

>   Pointers used to address the data in the packet. `head` points to the
>   beginning of the allocated space, `data` is the beginning of the
>   valid octets (and is usually slightly greater than `head`), `tail` is
>   the end of the valid octets, and `end` points to the maximum address
>   `tail` can reach. Another way to look at it is that the available
>   buffer space is `skb->end - skb->head`, and the currently used data
>   space is `skb->tail - skb->data`.

```
unsigned long len;
```

>   The length of the data itself (`skb->tail - skb->data`).

```
unsigned char ip_summed;
```

>   The checksum policy for this packet. The field is set by the driver
>   on incoming packets, as was described in "Packet Reception".

```
unsigned char pkt_type;
```

>   Packet classification used in delivering it. The driver is
>   responsible for setting it to `PACKET_HOST` (this packet is for me),
>   `PACKET_BROADCAST`, `PACKET_MULTICAST`, or `PACKET_OTHERHOST` (no, this packet is
>   not for me). Ethernet drivers don't modify `pkt_type` explicitly
>   because eth_type_trans does it for them.

The remaining fields in the structure are not particularly interesting.
They are used to maintain lists of buffers, to account for memory
belonging to the socket that owns the buffer, and so on.

Functions Acting on Socket Buffers

Network devices that use a `sock_buff` act on the structure by means of the official interface functions. Many functions operate on socket buffers; here are the most interesting ones:

```
struct sk_buff *alloc_skb(unsigned int len, int priority);
struct sk_buff *dev_alloc_skb(unsigned int len);
```

Allocate a buffer. The alloc_skb function allocates a buffer and initializes both `skb->data` and `skb->tail` to `skb->head`. The dev_alloc_skb function is a shortcut that calls alloc_skb with `GFP_ATOMIC` priority and reserves some space between `skb->head` and `skb->data`. This data space is used for optimizations within the network layer and should not be touched by the driver.

```
void kfree_skb(struct sk_buff *skb);
void dev_kfree_skb(struct sk_buff *skb);
```

Free a buffer. The kfree_skb call is used internally by the kernel. A driver should use dev_kfree_skb instead, which is intended to be safe to call from driver context.

```
unsigned char *skb_put(struct sk_buff *skb, int len);
unsigned char *__skb_put(struct sk_buff *skb, int len);
```

These inline functions update the `tail` and `len` fields of the `sk_buff` structure; they are used to add data to the end of the buffer. Each function's return value is the previous value of `skb->tail` (in other words, it points to the data space just created). Drivers can use the return value to copy data by invoking `ins(ioaddr, skb_put(...))` or `memcpy(skb_put(...), data, len)`. The difference between the two functions is that skb_put checks to be sure that the data will fit in the buffer, whereas __skb_put omits the check.

```
unsigned char *skb_push(struct sk_buff *skb, int len);
unsigned char *__skb_push(struct sk_buff *skb, int len);
```

These functions decrement `skb->data` and increment `skb->len`. They are similar to skb_put, except that data is added to the beginning of the packet instead of the end. The return value points to the data space just created. The functions are used to add a hardware header before transmitting a packet. Once again, __skb_push differs in that it does not check for adequate available space.

```
int skb_tailroom(struct sk_buff *skb);
```

This function returns the amount of space available for putting data in the buffer. If a driver puts more data into the buffer than it can hold, the system panics. Although you might object that a printk would be sufficient to tag the error, memory corruption is so harmful to the system that the developers decided to take definitive action. In practice, you shouldn't need to check the available space if the buffer has been correctly allocated. Since drivers usually get the packet size before allocating a buffer, only a severely broken driver will put too much data in the buffer,

and a panic might be seen as due punishment.

```
int skb_headroom(struct sk_buff *skb);
```

> Returns the amount of space available in front of `data`, that is,
> how many octets one can "push" to the buffer.

```
void skb_reserve(struct sk_buff *skb, int len);
```

> This function increments both `data` and `tail`. The function can be
> used to reserve headroom before filling the buffer. Most Ethernet
> interfaces reserve 2 bytes in front of the packet; thus, the IP
> header is aligned on a 16-byte boundary, after a 14-byte Ethernet
> header. snull does this as well, although the instruction was not
> shown in "Packet Reception" to avoid introducing extra concepts at
> that point.

```
unsigned char *skb_pull(struct sk_buff *skb, int len);
```

> Removes data from the head of the packet. The driver won't need to
> use this function, but it is included here for completeness. It
> decrements `skb->len` and increments `skb->data`; this is how the hardware
> header (Ethernet or equivalent) is stripped from the beginning of
> incoming packets.

The kernel defines several other functions that act on socket buffers,
but they are meant to be used in higher layers of networking code, and
the driver won't need them.

MAC Address Resolution

An interesting issue with Ethernet communication is how to associate the
MAC addresses (the interface's unique hardware ID) with the IP number.
Most protocols have a similar problem, but we concentrate on the
Ethernet-like case here. We'll try to offer a complete description of
the issue, so we will show three situations: ARP, Ethernet headers
without ARP (like plip), and non-Ethernet headers.

Using ARP with Ethernet

The usual way to deal with address resolution is by using ARP, the
Address Resolution Protocol. Fortunately, ARP is managed by the kernel,
and an Ethernet interface doesn't need to do anything special to support
ARP. As long as `dev->addr` and `dev->addr_len` are correctly assigned at open
time, the driver doesn't need to worry about resolving IP numbers to
physical addresses; ether_setup assigns the correct device methods to
`dev->hard_header` and `dev->rebuild_header`.

Although the kernel normally handles the details of address resolution
(and caching of the results), it calls upon the interface driver to help
in the building of the packet. After all, the driver knows about the
details of the physical layer header, while the authors of the
networking code have tried to insulate the rest of the kernel from that
knowledge. To this end, the kernel calls the driver's hard_header method

to lay out the packet with the results of the ARP query. Normally,
Ethernet driver writers need not know about this process -- the common
Ethernet code takes care of everything.

Overriding ARP

Simple point-to-point network interfaces such as plip might benefit from
using Ethernet headers, while avoiding the overhead of sending ARP
packets back and forth. The sample code in snull also falls into this
class of network devices. snullcannot use ARP because the driver changes
IP addresses in packets being transmitted, and ARP packets exchange IP
addresses as well. Although we could have implemented a simple ARP reply
generator with little trouble, it is more illustrative to show how to
handle physical-layer headers directly.

If your device wants to use the usual hardware header without running
ARP, you need to override the default dev->hard_header method. This is how
snull implements it, as a very short function.

```
int snull_header(struct sk_buff *skb, struct net_device *dev,
                 unsigned short type, void *daddr, void *saddr,
                 unsigned int len)
{
    struct ethhdr *eth = (struct ethhdr *)skb_push(skb,ETH_HLEN);

    eth->h_proto = htons(type);
    memcpy(eth->h_source, saddr ? saddr : dev->dev_addr, dev->addr_len);
    memcpy(eth->h_dest,   daddr ? daddr : dev->dev_addr, dev->addr_len);
    eth->h_dest[ETH_ALEN-1]    ^= 0x01;   /* dest is us xor 1 */
    return (dev->hard_header_len);
}
```

The function simply takes the information provided by the kernel and
formats it into a standard Ethernet header. It also toggles a bit in the
destination Ethernet address, for reasons described later.

When a packet is received by the interface, the hardware header is used
in a couple of ways by eth_type_trans. We have already seen this call in
snull_rx:

```
skb->protocol = eth_type_trans(skb, dev);
```

The function extracts the protocol identifier (ETH_P_IP in this case)
from the Ethernet header; it also assigns skb->mac.raw, removes the
hardware header from packet data (with skb_pull), and sets skb->pkt_type.
This last item defaults to PACKET_HOST at skb allocation (which indicates
that the packet is directed to this host), and eth_type_trans changes it
according to the Ethernet destination address. If that address does not
match the address of the interface that received it, the pkt_type field
will be set to PACKET_OTHERHOST. Subsequently, unless the interface is in
promiscuous mode, netif_rx will drop any packet of type PACKET_OTHERHOST.
For this reason, snull_header is careful to make the destination
hardware address match that of the "receiving" interface.

If your interface is a point-to-point link, you won't want to receive
unexpected multicast packets. To avoid this problem, remember that a

destination address whose first octet has 0 as the least significant bit (LSB) is directed to a single host (i.e., it is either `PACKET_HOST` or `PACKET_OTHERHOST`). The plip driver uses 0xfc as the first octet of its hardware address, while snull uses 0x00. Both addresses result in a working Ethernet-like point-to-point link.

Non-Ethernet Headers

We have just seen that the hardware header contains some information in addition to the destination address, the most important being the communication protocol. We now describe how hardware headers can be used to encapsulate relevant information. If you need to know the details, you can extract them from the kernel sources or the technical documentation for the particular transmission medium. Most driver writers will be able to ignore this discussion and just use the Ethernet implementation.

It's worth noting that not all information has to be provided by every protocol. A point-to-point link such as plipor snull could avoid transferring the whole Ethernet header without losing generality. The hard_header device method, shown earlier as implemented by snull_header, receives the delivery information -- both protocol-level and hardware addresses -- from the kernel. It also receives the 16-bit protocol number in the `type` argument; IP, for example, is identified by `ETH_P_IP`. The driver is expected to correctly deliver both the packet data and the protocol number to the receiving host. A point-to-point link could omit addresses from its hardware header, transferring only the protocol number, because delivery is guaranteed independent of the source and destination addresses. An IP-only link could even avoid transmitting any hardware header whatsoever.

When the packet is picked up at the other end of the link, the receiving function in the driver should correctly set the fields `skb->protocol`, `skb->pkt_type`, and `skb->mac.raw`.

`skb->mac.raw` is a char pointer used by the address-resolution mechanism implemented in higher layers of the networking code (for instance, net/ipv4/arp.c). It must point to a machine address that matches `dev->type`. The possible values for the device type are defined in <linux/if_arp.h>; Ethernet interfaces use `ARPHRD_ETHER`. For example, here is how eth_type_trans deals with the Ethernet header for received packets:

```
skb->mac.raw = skb->data;
skb_pull(skb, dev->hard_header_len);
```

In the simplest case (a point-to-point link with no headers), `skb->mac.raw` can point to a static buffer containing the hardware address of this interface, `protocol` can be set to `ETH_P_IP`, and `packet_type` can be left with its default value of `PACKET_HOST`.

Because every hardware type is unique, it is hard to give more specific advice than already discussed. The kernel is full of examples, however. See, for example, the AppleTalk driver (drivers/net/appletalk/cops.c), the infrared drivers (such as drivers/net/irda/smc_ircc.c), or the PPP driver (drivers/net/ppp_generic.c).

Custom ioctl Commands

We have seen that the ioctl system call is implemented for sockets; SIOCSIFADDR and SIOCSIFMAP are examples of "socket ioctls." Now let's see how the third argument of the system call is used by networking code.

When the ioctl system call is invoked on a socket, the command number is one of the symbols defined in <linux/sockios.h>, and the function sock_ioctl directly invokes a protocol-specific function (where "protocol" refers to the main network protocol being used, for example, IP or AppleTalk).

Any ioctl command that is not recognized by the protocol layer is passed to the device layer. These device-related ioctl commands accept a third argument from user space, a struct ifreq *. This structure is defined in <linux/if.h>. The SIOCSIFADDR and SIOCSIFMAP commands actually work on the ifreq structure. The extra argument to SIOCSIFMAP, although defined as ifmap, is just a field of ifreq.

In addition to using the standardized calls, each interface can define its own ioctl commands. The plip interface, for example, allows the interface to modify its internal timeout values via ioctl. The ioctlimplementation for sockets recognizes 16 commands as private to the interface: SIOCDEVPRIVATE through SIOCDEVPRIVATE+15.

When one of these commands is recognized, dev->do_ioctl is called in the relevant interface driver. The function receives the same struct ifreq * pointer that the general-purpose ioctl function uses:

        int (*do_ioctl)(struct net_device *dev, struct ifreq *ifr, int cmd);

The ifr pointer points to a kernel-space address that holds a copy of the structure passed by the user. After do_ioctl returns, the structure is copied back to user space; the driver can thus use the private commands to both receive and return data.

The device-specific commands can choose to use the fields in struct ifreq, but they already convey a standardized meaning, and it's unlikely that the driver can adapt the structure to its needs. The field ifr_data is a caddr_t item (a pointer) that is meant to be used for device-specific needs. The driver and the program used to invoke its ioctl commands should agree about the use of ifr_data. For example, pppstats uses device-specific commands to retrieve information from the ppp interface driver.

It's not worth showing an implementation of do_ioctl here, but with the information in this chapter and the kernel examples, you should be able to write one when you need it. Note, however, that the plipimplementation uses ifr_data incorrectly and should not be used as an example for an ioctlimplementation.

Statistical Information

The last method a driver needs is get_stats. This method returns a

pointer to the statistics for the device. Its implementation is pretty easy; the one shown works even when several interfaces are managed by the same driver, because the statistics are hosted within the device data structure.

```
struct net_device_stats *snull_stats(struct net_device *dev)
{
    struct snull_priv *priv = (struct snull_priv *) dev->priv;
    return &priv->stats;
}
```

The real work needed to return meaningful statistics is distributed throughout the driver, where the various fields are updated. The following list shows the most interesting fields in struct net_device_stats.

unsigned long rx_packets;
unsigned long tx_packets;

These fields hold the total number of incoming and outgoing packets successfully transferred by the interface.

unsigned long rx_bytes;
unsigned long tx_bytes;

The number of bytes received and transmitted by the interface. These fields were added in the 2.2 kernel.

unsigned long rx_errors;
unsigned long tx_errors;

The number of erroneous receptions and transmissions. There's no end of things that can go wrong with packet transmission, and the net_device_stats structure includes six counters for specific receive errors and five for transmit errors. See <linux/netdevice.h> for the full list. If possible, your driver should maintain detailed error statistics, because they can be most helpful to system administrators trying to track down a problem.

unsigned long rx_dropped;
unsigned long tx_dropped;

The number of packets dropped during reception and transmission. Packets are dropped when there's no memory available for packet data. tx_dropped is rarely used.

unsigned long collisions;

The number of collisions due to congestion on the medium.

unsigned long multicast;

The number of multicast packets received.

It is worth repeating that the get_stats method can be called at any time -- even when the interface is down -- so the driver should not release statistic information when running the stop method.

Multicasting

A multicast packet is a network packet meant to be received by more than
one host, but not by all hosts. This functionality is obtained by
assigning special hardware addresses to groups of hosts. Packets
directed to one of the special addresses should be received by all the
hosts in that group. In the case of Ethernet, a multicast address has
the least significant bit of the first address octet set in the
destination address, while every device board has that bit clear in its
own hardware address.

The tricky part of dealing with host groups and hardware addresses is
performed by applications and the kernel, and the interface driver
doesn't need to deal with these problems.

Transmission of multicast packets is a simple problem because they look
exactly like any other packets. The interface transmits them over the
communication medium without looking at the destination address. It's
the kernel that has to assign a correct hardware destination address;
the hard_header device method, if defined, doesn't need to look in the
data it arranges.

The kernel handles the job of tracking which multicast addresses are of
interest at any given time. The list can change frequently, since it is
a function of the applications that are running at any given time and
the user's interest. It is the driver's job to accept the list of
interesting multicast addresses and deliver to the kernel any packets
sent to those addresses. How the driver implements the multicast list is
somewhat dependent on how the underlying hardware works. Typically,
hardware belongs to one of three classes, as far as multicast is
concerned:

- Interfaces that cannot deal with multicast. These interfaces either
  receive packets directed specifically to their hardware address
  (plus broadcast packets), or they receive every packet. They can
  receive multicast packets only by receiving every packet, thus
  potentially overwhelming the operating system with a huge number of
  "uninteresting" packets. You don't usually count these interfaces
  as multicast capable, and the driver won't set IFF_MULTICAST in dev-
  >flags.

  Point-to-point interfaces are a special case, because they always
  receive every packet without performing any hardware filtering.

- Interfaces that can tell multicast packets from other packets
  (host-to-host or broadcast). These interfaces can be instructed to
  receive every multicast packet and let the software determine if
  this host is a valid recipient. The overhead introduced in this
  case is acceptable, because the number of multicast packets on a
  typical network is low.

- Interfaces that can perform hardware detection of multicast
  addresses. These interfaces can be passed a list of multicast
  addresses for which packets are to be received, and they will

ignore other multicast packets. This is the optimum case for the
kernel, because it doesn't waste processor time dropping
"uninteresting" packets received by the interface.

The kernel tries to exploit the capabilities of high-level interfaces by
supporting at its best the third device class, which is the most
versatile. Therefore, the kernel notifies the driver whenever the list
of valid multicast addresses is changed, and it passes the new list to
the driver so it can update the hardware filter according to the new
information.

Kernel Support for Multicasting

Support for multicast packets is made up of several items: a device
method, a data structure and device flags.

`void (*dev->set_multicast_list)(struct net_device *dev);`

>   This device method is called whenever the list of machine addresses
>   associated with the device changes. It is also called when `dev->flags`
>   is modified, because some flags (e.g., `IFF_PROMISC`) may also require
>   you to reprogram the hardware filter. The method receives a pointer
>   to `struct net_device` as an argument and returns `void`. A driver not
>   interested in implementing this method can leave the field set to
>   `NULL`.

`struct dev_mc_list *dev->mc_list;`

>   This is a linked list of all the multicast addresses associated
>   with the device. The actual definition of the structure is
>   introduced at the end of this section.

`int dev->mc_count;`

>   The number of items in the linked list. This information is
>   somewhat redundant, but checking `mc_count` against 0 is a useful
>   shortcut for checking the list.

`IFF_MULTICAST`

>   Unless the driver sets this flag in `dev->flags`, the interface won't
>   be asked to handle multicast packets. The set_multicast_list method
>   will nonetheless be called when `dev->flags` changes, because the
>   multicast list may have changed while the interface was not active.

`IFF_ALLMULTI`

>   This flag is set in `dev->flags` by the networking software to tell the
>   driver to retrieve all multicast packets from the network. This
>   happens when multicast routing is enabled. If the flag is set, `dev->mc_list` shouldn't be used to filter multicast packets.

`IFF_PROMISC`

>   This flag is set in `dev->flags` when the interface is put into
>   promiscuous mode. Every packet should be received by the interface,

independent of `dev->mc_list`.

The last bit of information needed by the driver developer is the
definition of `struct dev_mc_list`, which lives in `<linux/netdevice.h>`.

```
struct dev_mc_list {
    struct dev_mc_list  *next;            /* Next address in list */
    __u8                dmi_addr[MAX_ADDR_LEN]; /* Hardware address */
    unsigned char       dmi_addrlen;    /* Address length */
    int                 dmi_users;      /* Number of users */
    int                 dmi_gusers;     /* Number of groups */
};
```

Because multicasting and hardware addresses are independent of the
actual transmission of packets, this structure is portable across
network implementations, and each address is identified by a string of
octets and a length, just like `dev->dev_addr`.

A Typical Implementation

The best way to describe the design of set_multicast_list is to show you
some pseudocode.

The following function is a typical implementation of the function in a
full-featured (`ff`) driver. The driver is full featured in that the
interface it controls has a complex hardware packet filter, which can
hold a table of multicast addresses to be received by this host. The
maximum size of the table is `FF_TABLE_SIZE`.

All the functions prefixed with `ff_` are placeholders for hardware-
specific operations.

```
void ff_set_multicast_list(struct net_device *dev)
{
    struct dev_mc_list *mcptr;

    if (dev->flags & IFF_PROMISC) {
        ff_get_all_packets();
        return;
    }
    /* If there's more addresses than we handle, get all multicast
    packets and sort them out in software. */
    if (dev->flags & IFF_ALLMULTI || dev->mc_count > FF_TABLE_SIZE) {
        ff_get_all_multicast_packets();
        return;
    }
    /* No multicast?  Just get our own stuff */
    if (dev->mc_count == 0) {
        ff_get_only_own_packets();
        return;
    }
    /* Store all of the multicast addresses in the hardware filter */
    ff_clear_mc_list();
    for (mc_ptr = dev->mc_list; mc_ptr; mc_ptr = mc_ptr->next)
        ff_store_mc_address(mc_ptr->dmi_addr);
    ff_get_packets_in_multicast_list();
}
```

This implementation can be simplified if the interface cannot store a
multicast table in the hardware filter for incoming packets. In that

case, `FF_TABLE_SIZE` reduces to 0 and the last four lines of code are not needed.

As was mentioned earlier, even interfaces that can't deal with multicast packets need to implement the set_multicast_list method to be notified about changes in `dev->flags`. This approach could be called a "nonfeatured" (`nf`) implementation. The implementation is very simple, as shown by the following code:

```
void nf_set_multicast_list(struct net_device *dev)
{
    if (dev->flags & IFF_PROMISC)
        nf_get_all_packets();
    else
        nf_get_only_own_packets();
}
```

Implementing `IFF_PROMISC` is important, because otherwise the user won't be able to run tcpdump or any other network analyzers. If the interface runs a point-to-point link, on the other hand, there's no need to implement set_multicast_list at all, because users receive every packet anyway.

Backward Compatibility

Version 2.3.43 of the kernel saw a major rework of the networking subsystem. The new "softnet" implementation was a great improvement in terms of performance and clean design. It also, of course, brought changes to the network driver interface -- though fewer than one might have expected.

Differences in Linux 2.2

First of all, Linux 2.3.14 renamed the network device structure, which had always been `struct device`, to `struct net_device`. The new name is certainly more appropriate, since the structure was never meant to describe devices in general.

Prior to version 2.3.43, the functions netif_start_queue, netif_stop_queue, and netif_wake_queue did not exist. Packet transmission was, instead, controlled by three fields in the `device` structure, and sysdep.himplements the three functions using the three fields when compiling for 2.2 or 2.0.

`unsigned char start;`

> This variable indicated that the interface was ready for operations; it was normally set to 1 in the driver's openmethod. The current implementation is to call netif_start_queue instead.

`unsigned long interrupt;`

> `interrupt` was used to indicate that the device was servicing an interrupt -- accordingly, it was set to 1 at the beginning of the interrupt handler and to 0 before returning. It was never a

substitute for proper locking, and its use has been replaced with internal spinlocks.

`unsigned long tbusy;`

> When nonzero, this variable indicated that the device could handle no more outgoing packets. Where a 2.4 driver will call netif_stop_queue, older drivers would set `tbusy` to 1. Restarting the queue required setting `tbusy` back to 0 and calling `mark_bh(NET_BH)`.

Normally, setting `tbusy` was sufficient to ensure that the driver's hard_start_xmit method would not be called. However, if the networking system decided that a transmitter lockup must have occurred, it would call that method anyway. There was no tx_timeout method before softnet was integrated. Thus, pre-softnet drivers had to explicitly check for a call to hard_start_xmit when `tbusy` was set and react accordingly.

The type of the `name` field in `struct device` was different. The 2.2 version was simply

`char *name;`

Thus, the storage for the interface name had to be allocated separately, and `name` assigned to point to that storage. Usually the device name was stored in a static variable within the driver. The `%d` notation for dynamically assigned interface names was not present in 2.2; instead, if the name began with a null byte or a space character, the kernel would allocate the next `eth` name. The 2.4 kernel still implements this behavior, but its use is deprecated. Starting with 2.5, only the `%d` format is likely to be recognized.

The `owner` field (and the `SET_MODULE_OWNER` macro) were added in kernel 2.4.0-test11, just before the official stable release. Previously, network driver modules had to maintain their own use counts. sysdep.h defines an empty `SET_MODULE_OWNER` for kernels that do not have it; portable code should also continue to manage its use count manually (in addition to letting the networking system do it).

The link state functions (netif_carrier_on and netif_carrier_off) did not exist in the 2.2 kernel. The kernel simply did without that information in those days.

Further Differences in Linux 2.0

The 2.1 development series also saw its share of changes to the network driver interface. Most took the form of small changes to function prototypes, rather than sweeping changes to the network code as a whole.

Interface statistics were kept in a structure called `struct lenet_statistics`, defined in `<linux/if_ether.h>`. Even non-Ethernet drivers used this structure. The field names were all the same as the current `struct net_device_stats`, but the `rx_bytes` and `tx_bytes` fields were not present.

The 2.0 kernel handled transmitter lockups in the same way as 2.2 did. There was, however, an additional function:

```
void dev_tint(struct device *dev);
```

This function would be called by the driver after a lockup had been
cleared to restart the transmission of packets.

A couple of functions had different prototypes. dev_kfree_skb had a
second, integer argument that was either FREE_READ for incoming packets
(i.e., skbs allocated by the driver) or FREE_WRITE for outgoing packets
(skbs allocated by the networking code). Almost all calls to
dev_kfree_skb in network driver code used FREE_WRITE. The nonchecking
versions of the skb functions (such as __skb_push) did not exist;
sysdep.h in the sample code provides emulation for these functions under
2.0.

The rebuild_header method had a different set of arguments:

```
int (*rebuild_header) (void *eth, struct device *dev,
unsigned long raddr, struct sk_buff *skb);
```

The Linux kernel also made heavier use of rebuild_header; it did most of
the work that hard_header does now. When snull is compiled under Linux
2.0, it builds hardware headers as follows:

```
int snull_rebuild_header(void *buff, struct net_device *dev, unsigned long dst,
                struct sk_buff *skb)
{
    struct ethhdr *eth = (struct ethhdr *)buff;

    memcpy(eth->h_source, dev->dev_addr, dev->addr_len);
    memcpy(eth->h_dest, dev->dev_addr, dev->addr_len);
    eth->h_dest[ETH_ALEN-1]   ^= 0x01;   /* dest is us xor 1 */
    return 0;
}
```

The device methods for header caching were also significantly different
in this kernel. If your driver needs to implement these functions
directly (very few do), and it also needs to work with the 2.0 kernel,
see the definitions in <linux/netdevice.h> to see how things were done in
those days.

Probing and HAVE_DEVLIST

If you look at the source for almost any network driver in the kernel,
you will find some boilerplate that looks like this:

```
#ifdef HAVE_DEVLIST
/*
 * Support for an alternate probe manager,
 * which will eliminate the boilerplate below.
 */
struct netdev_entry netcard_drv =
{cardname, netcard_probe1, NETCARD_IO_EXTENT, netcard_portlist};
#else
/* Regular probe routine defined here */
```

Interestingly, this code has been around since the 1.1 development
series, but we are still waiting for the promised alternate probe

manager. It is probably safe to not worry about being prepared for this great change, especially since ideas for how to implement it will likely have changed in the intervening years.

Quick Reference

This section provides a reference for the concepts introduced in this chapter. It also explains the role of each header file that a driver needs to include. The lists of fields in the net_device and sk_buff structures, however, are not repeated here.

#include <linux/netdevice.h>

> This header hosts the definitions of struct net_device and struct net_device_stats, and includes a few other headers that are needed by network drivers.

int register_netdev(struct net_device *dev);
void unregister_netdev(struct net_device *dev);

> Register and unregister a network device.

SET_MODULE_OWNER(struct net_device *dev);

> This macro will store a pointer to the current module in the device structure (or in any structure with an owner field, actually); it is used to enable the networking subsystem to manage the module's use count.

netif_start_queue(struct net_device *dev);
netif_stop_queue(struct net_device *dev);
netif_wake_queue(struct net_device *dev);

> These functions control the passing of packets to the driver for transmission. No packets will be transmitted until netif_start_queue has been called. netif_stop_queue suspends transmission, and netif_wake_queue restarts the queue and pokes the network layer to restart transmitting packets.

void netif_rx(struct sk_buff *skb);

> This function can be called (including at interrupt time) to notify the kernel that a packet has been received and encapsulated into a socket buffer.

#include <linux/if.h>

> Included by netdevice.h, this file declares the interface flags (IFF_ macros) and struct ifmap, which has a major role in the ioctl implementation for network drivers.

void netif_carrier_off(struct net_device *dev);
void netif_carrier_on(struct net_device *dev);
int netif_carrier_ok(struct net_device *dev);

> The first two functions may be used to tell the kernel whether a carrier signal is currently present on the given interface.

netif_carrier_ok will test the carrier state as reflected in the device structure.

```
#include <linux/if_ether.h>
ETH_ALEN
ETH_P_IP
struct ethhdr;
```

Included by netdevice.h, if_ether.h defines all the ETH_ macros used to represent octet lengths (such as the address length) and network protocols (such as IP). It also defines the ethhdr structure.

```
#include <linux/skbuff.h>
```

The definition of struct sk_buff and related structures, as well as several inline functions to act on the buffers. This header is included by netdevice.h.

```
struct sk_buff *alloc_skb(unsigned int len, int priority);
struct sk_buff *dev_alloc_skb(unsigned int len);
void kfree_skb(struct sk_buff *skb);
void dev_kfree_skb(struct sk_buff *skb);
```

These functions handle the allocation and freeing of socket buffers. Drivers should normally use the dev_ variants, which are intended for that purpose.

```
unsigned char *skb_put(struct sk_buff *skb, int len);
unsigned char *__skb_put(struct sk_buff *skb, int len);
unsigned char *skb_push(struct sk_buff *skb, int len);
unsigned char *__skb_push(struct sk_buff *skb, int len);
```

These functions add data to an skb; skb_put puts the data at the end of the skb, while skb_push puts it at the beginning. The regular versions perform checking to ensure that adequate space is available; double-underscore versions leave those tests out.

```
int skb_headroom(struct sk_buff *skb);
int skb_tailroom(struct sk_buff *skb);
void skb_reserve(struct sk_buff *skb, int len);
```

These functions perform management of space within an skb. skb_headroom and skb_tailroom tell how much space is available at the beginning and end, respectively, of an skb. skb_reserve may be used to reserve space at the beginning of an skb, which must be empty.

```
unsigned char *skb_pull(struct sk_buff *skb, int len);
```

skb_pull will "remove" data from an skb by adjusting the internal pointers.

```
#include <linux/etherdevice.h>
void ether_setup(struct net_device *dev);
```

This function sets most device methods to the general-purpose implementation for Ethernet drivers. It also sets dev->flags and assigns the next available ethx name to dev->name if the first

character in the name is a blank space or the null character.

```
unsigned short eth_type_trans(struct sk_buff *skb, struct net_device *dev);
```

When an Ethernet interface receives a packet, this function can be called to set `skb->pkt_type`. The return value is a protocol number that is usually stored in `skb->protocol`.

```
#include <linux/sockios.h>
SIOCDEVPRIVATE
```

This is the first of 16 ioctl commands that can be implemented by each driver for its own private use. All the network ioctl commands are defined in sockios.h.

Back to: [Table of Contents](#)

Back to: [Linux Device Drivers, 2nd Edition](#)